

# Let's Evaluate

## Organization-Wide Product Alignment & Acceptance Document

| Field         | Value                                          |
|---------------|------------------------------------------------|
| Product       | Let's Evaluate                                 |
| Document type | Organization-wide acceptance & alignment brief |
| Version       | 1.0                                            |
| Status        | Draft for team review                          |
| Audience      | TA team, interviewers, engineering, leadership |
| Date          | July 2026                                      |

---

## Table of Contents

1. Executive Summary
2. Problem Statement
3. Vision & Objectives
4. Users & Personas
5. Product Scope — Feature Inventory
6. Target Workflow
7. Gap Analysis & Recommendations
8. Technical Architecture
9. Phased Delivery Plan
10. Success Metrics
11. Risks & Mitigations
12. Team Acceptance Checklist
13. Appendix — Codebase Mapping

---

## 1. Executive Summary

**Let's Evaluate** is an end-to-end technical hiring platform that keeps recruiters, interviewers, and (later) administrators in a single portal — from resume screening through structured interviews to candidate communication and reporting.

**Core promise:** Reduce wasted recruiter and interviewer time by aligning project requirements, role expectations, interview content, and outcomes *\*before\** and *\*during\** interviews — not after.

**Strategic direction:** Build and validate internally first. Once proven to save time and improve hiring quality, package as a white-label SaaS product (similar to Zoho) with per-organization branding, domains, and configuration.

**Current state:** A working prototype exists (Streamlit/Python) and an active Next.js rewrite with organization tenancy, role-based access, white-label branding, TA→interviewer workflow, and AI-powered resume analysis. Several features from this ideation are partially built; others are planned but not yet implemented.

---

## 2. Problem Statement

### 2.1 Current Pain Points

| Pain Area                          | Impact                                                                                 |
|------------------------------------|----------------------------------------------------------------------------------------|
| Misaligned screening               | TA screens without full project context → unsuitable candidates reach the panel        |
| Panel mismatch                     | Interviewers assigned without clear role/project fit                                   |
| High rejection rate post-interview | Time lost for recruiters, interviewers, and candidates                                 |
| Fragmented tooling                 | Resume review, evaluation, assignment, and email happen across multiple tools          |
| Template mismatch                  | Emails describe Lead/Developer/Architect roles inconsistently vs. the actual interview |
| No single source of truth          | No shared timeline, status, or audit trail across the hiring flow                      |
| Slow closure                       | Long cycles from screen → interview → decision → candidate communication               |

### 2.2 Root Causes

1. Requirements (project + role) are not enforced at screening time.
  2. Interview preparation is ad hoc — no structured question bank or draft workflow.
  3. Communication is disconnected from evaluation state.
  4. No visibility into panel availability, ETAs, or pipeline health.
- 

## 3. Vision & Objectives

### 3.1 Vision

One configurable portal where hiring teams run the full technical evaluation lifecycle — with AI assistance where it saves time, and human judgment where it matters.

### 3.2 Objectives

| #  | Objective                    | Success Indicator                                                    |
|----|------------------------------|----------------------------------------------------------------------|
| O1 | <b>Unified workflow</b>      | No context switching for resume, evaluation, assignment, and mail    |
| O2 | <b>Requirement alignment</b> | Screening and interview content tied to project + role configuration |
| O3 | <b>Time savings</b>          | Measurable reduction in unsuitable interviews and cycle time         |
| O4 | <b>Configurable product</b>  | Branding, domain, templates, and flows configurable per deployment   |
| O5 | <b>Auditability</b>          | Full candidate timeline and decision history                         |
| O6 | <b>Commercial readiness</b>  | Multi-tenant / resale path documented and validated                  |

### 3.3 Non-Goals (Initial Validation Phase)

- Full ATS replacement (Greenhouse, Lever, etc.)
- Built-in video conferencing (integrate instead)
- Fully autonomous hiring decisions without human review
- Multi-org SaaS billing in MVP (single org per deployment first)

## 4. Users & Personas

| Persona                     | Primary Needs                                                                           | Module              |
|-----------------------------|-----------------------------------------------------------------------------------------|---------------------|
| <b>Recruiter (TA)</b>       | Pipeline visibility, project/role setup, screening, assignment, candidate communication | Recruiter           |
| <b>Interviewer (Panel)</b>  | Schedule visibility, preparation, structured evaluation, submit report                  | Interviewer         |
| <b>Admin</b> (future)       | Users, org settings, projects, roles, policies                                          | Admin               |
| <b>Candidate</b> (indirect) | Clear, accurate communication; timely updates                                           | Email triggers only |

### Role-Based Access (Target)

| Role        | Permissions                                                |
|-------------|------------------------------------------------------------|
| ta          | Full recruiter workflow; cannot override org configuration |
| interviewer | Assigned cases only; question bank; submit evaluations     |
| admin       | Users, org config, mail templates, hierarchy, reporting    |

\*Status: admin, ta, and interviewer RBAC already implemented in the Next.js application.\*

## 5. Product Scope — Feature Inventory

### 5.1 Platform & Authentication

| Feature                   | Description                          | Priority | Status                               |
|---------------------------|--------------------------------------|----------|--------------------------------------|
| Sign up / Sign in         | Domain-restricted email registration | P0       | <b>Built</b>                         |
| Two-factor authentication | TOTP or email OTP                    | P1       | <b>Not built</b>                     |
| Org-wide SSO              | Microsoft Entra ID / Azure AD        | P1       | <b>Partial</b> — provider scaffolded |
| Session security          | JWT, bcrypt, rate limits             | P0       | <b>Built</b>                         |
| White-label branding      | Name, tagline, colors, logo, favicon | P0       | <b>Built</b>                         |

### 5.2 Recruiter Module

| #  | Feature                                                 | Priority | Status           | Notes                                       |
|----|---------------------------------------------------------|----------|------------------|---------------------------------------------|
| R1 | Dashboard — ongoing interviews, panel availability, ETA | P0       | <b>Partial</b>   | Stats exist; panel calendar and ETA missing |
| R2 | Project configuration                                   | P0       | <b>Built</b>     | Tech stack, CRUD                            |
| R3 | Role configuration                                      | P0       | <b>Built</b>     | Requirements per role                       |
| R4 | Panel timeline & slot booking                           | P1       | <b>Not built</b> | Critical for scheduling alignment           |

| #   | Feature                                               | Priority | Status           | Notes                                         |
|-----|-------------------------------------------------------|----------|------------------|-----------------------------------------------|
| R5  | Mail templates (schedule, receive, shortlist, reject) | P0       | <b>Not built</b> | Required for candidate communication          |
| R6  | Candidate-triggered status emails                     | P1       | <b>Not built</b> | Configurable deployment triggers              |
| R7  | Interview hierarchy flow configuration                | P1       | <b>Partial</b>   | Status enum exists; UI-driven stages pending  |
| R8  | AI resume evaluation + cross-project suggestion       | P0       | <b>Partial</b>   | Analysis built; cross-project routing pending |
| R9  | Assign interviewer                                    | P0       | <b>Built</b>     | Assignment workflow exists                    |
| R10 | Pipeline / kanban view                                | P1       | <b>Built</b>     | Needs main navigation placement               |

### 5.3 Interviewer Module

| #  | Feature                                                          | Priority | Status            | Notes                          |
|----|------------------------------------------------------------------|----------|-------------------|--------------------------------|
| I1 | Dashboard — schedules & assignments                              | P0       | <b>Partial</b>    | Assignment list exists         |
| I2 | View TA screening report (reuse AI output)                       | P0       | <b>Built</b>      | Saves duplicate AI cost        |
| I3 | Question generation (resume, role, scenario, coding, analytical) | P0       | <b>Partial</b>    | Standard + resume-based built  |
| I4 | Custom prompt for questions                                      | P1       | <b>Not in web</b> | Exists in Streamlit legacy app |
| I5 | Question bank with categories                                    | P1       | <b>Partial</b>    | CRUD exists; taxonomy pending  |
| I6 | Per-question rating + comments                                   | P0       | <b>Partial</b>    | Review schema exists           |
| I7 | Draft / prepare before interview                                 | P0       | <b>Built</b>      | Draft save and resume          |

| #   | Feature                              | Priority | Status         | Notes                                 |
|-----|--------------------------------------|----------|----------------|---------------------------------------|
| I8  | AI-refined final evaluation comments | P1       | <b>Partial</b> | Code exists; not wired to UI          |
| I9  | Submit → PDF report to recruiter     | P0       | <b>Partial</b> | Streamlit has PDF; web is placeholder |
| I10 | Close case on interviewer side       | P0       | <b>Built</b>   | Assignment completion flow            |

## 5.4 Reporting (Both Parties)

| Feature                       | Priority | Status                            |
|-------------------------------|----------|-----------------------------------|
| Interview counts by status    | P1       | <b>Partial</b> — basic stats only |
| Visual charts (org dashboard) | P2       | <b>Not built</b>                  |
| Exportable reports            | P2       | <b>Not built</b>                  |

## 5.5 Admin Module (Future)

| Feature                                           | Phase   |
|---------------------------------------------------|---------|
| User management (invite, deactivate, role change) | Phase 2 |
| Org-wide settings UI (replace env-only config)    | Phase 2 |
| Audit log viewer                                  | Phase 2 |
| Template & flow management UI                     | Phase 2 |

# 6. Target Workflow

## 6.1 End-to-End Flow

```

TA uploads resume
  → AI resume analysis (vs. project + role)
    → TA decision: Proceed / Hold / Reject
      → [Reject] Send reject template email
      → [Hold] Hold with notes
      → [Proceed] Assign interviewer + book slot
        → Interviewer views screening report (no re-analysis)
          → Generate or reuse questions from bank
            → Optional: save draft and prepare
              → Conduct interview (rate each question + comments)
                → AI refine final evaluation comments
                  → Submit PDF report to recruiter
                    → Recruiter reviews + sends candidate email
                      → Archive + reporting metrics

```

[If mismatch] AI suggests alternate project → TA confirms reassignment

## 6.2 Candidate Status Lifecycle

draft → screening → screened\_hold / screened\_rejected  
→ ready\_for\_interview → assigned → interview\_in\_progress  
→ interview\_complete → selected / rejected / hold

\*Status: Candidate status enum already defined in the database schema.\*

---

## 7. Gap Analysis & Recommendations

### 7.1 Strengths of Current Ideation

- **Single-portal workflow** — primary differentiator and correct focus
- **Reuse TA AI analysis for interviewers** — smart cost control
- **Draft preparation before live interview** — reduces interviewer stress
- **Configurable branding for resale** — already started in codebase
- **Structured question bank with categories** — reduces repeat AI calls
- **Interview hierarchy configuration** — supports Lead / Developer / Architect flows

### 7.2 Scope Refinements

| Feature                             | Recommendation                                                                             |
|-------------------------------------|--------------------------------------------------------------------------------------------|
| <b>Panel timeline &amp; booking</b> | Phase 1: manual availability windows + slot pick.<br>Phase 2: Outlook/Google calendar sync |
| <b>Candidate mail triggers</b>      | Start with recruiter-initiated + status-change auto-send; add reply parsing later          |
| <b>Cross-project AI routing</b>     | Suggestions only — TA must confirm reassignment                                            |

### 7.3 Recommended Additions

| Gap                             | Why It Matters                                          | Priority |
|---------------------------------|---------------------------------------------------------|----------|
| <b>Candidate record</b>         | Name, email, phone, source, consent                     | P0       |
| <b>Interview calendar / ICS</b> | Send calendar invites with correct role context         | P1       |
| <b>Notification system</b>      | In-app + email when assigned, submitted, overdue        | P1       |
| <b>SLA / ETA tracking</b>       | Dashboard promises ETAs — needs due dates and reminders | P1       |

| Gap                                  | Why It Matters                                                      | Priority |
|--------------------------------------|---------------------------------------------------------------------|----------|
| <b>Template variables</b>            | {{candidate_name}},<br>{{role}}, {{project}},<br>{{interview_date}} | P0       |
| <b>Evaluation rubric / scorecard</b> | Weighted criteria per role level<br>(Lead vs. Architect)            | P1       |
| <b>Bias &amp; compliance</b>         | PII handling, data retention,<br>export/delete                      | P1       |
| <b>Activity audit log</b>            | Who changed what, when                                              | P1       |
| <b>Bulk import/export</b>            | Projects, roles, questions                                          | P2       |
| <b>Legal pages</b>                   | Privacy policy, terms and<br>conditions                             | P2       |
| <b>Forgot password</b>               | Self-service reset for web app                                      | P1       |
| <b>AI cost controls</b>              | Per-org budgets, caching, question<br>bank reuse metrics            | P1       |
| <b>Multi-interviewer rounds</b>      | Round 1 screen → Round 2 tech<br>→ Round 3 architecture             | P2       |

#### 7.4 AI Model Strategy

| Use Case                                        | Recommended Model                   | Rationale                                                  |
|-------------------------------------------------|-------------------------------------|------------------------------------------------------------|
| Resume analysis (scoring, tech match, concerns) | <b>GPT-4o</b> or <b>GPT-4.1</b>     | Structured reasoning; higher accuracy for hiring decisions |
| Cross-project fit suggestion                    | Same as above + project index in DB | One additional prompt with all active projects             |
| Question generation (standard, resume-based)    | <b>GPT-4o-mini</b>                  | Fast, cost-effective for question lists                    |
| Coding / scenario questions                     | <b>GPT-4o</b>                       | Stronger reasoning for technical scenarios                 |
| Evaluation comment refinement                   | <b>GPT-4o-mini</b>                  | Summarization task; code already exists                    |
| Question deduplication (future)                 | <b>text-embedding-3-small</b>       | Avoid regenerating similar questions                       |

#### 7.5 AI Cost Control Strategy

1. Run full resume AI **once** at TA screening — interviewers read cached result.
2. Check question bank before calling AI.
3. Cache analysis by resume hash + project + role version.




- 4. Rate limit per user (currently: 30 requests per window).
- 5. Track token usage per organization in audit or dedicated usage table.

## 8. Technical Architecture

### 8.1 Technology Stack

| Layer          | Technology                                        | Notes                              |
|----------------|---------------------------------------------------|------------------------------------|
| Frontend       | Next.js 16, React 19, TypeScript                  | App Router                         |
| Styling        | Tailwind CSS 4, design tokens                     | White-label via CSS variables      |
| API            | Next.js API routes / Server Actions               | No separate backend for MVP        |
| Database       | PostgreSQL                                        | Neon (dev) → Azure Postgres (prod) |
| ORM            | Drizzle ORM                                       | Schema-driven migrations           |
| Auth           | Auth.js (next-auth v5)                            | Credentials + Microsoft Entra ID   |
| File storage   | Local dev / S3-compatible                         | Resumes, generated PDFs            |
| Email          | Resend, SendGrid, or Azure Communication Services | Transactional templates            |
| PDF generation | @react-pdf/renderer or Puppeteer                  | Server-side report generation      |
| AI             | OpenAI API (direct SDK)                           | No LangChain in web stack          |
| Hosting        | Vercel (app) + Neon/Azure (DB) + S3 (files)       | Documented deployment path         |
| Observability  | Sentry, Vercel Analytics                          | Add before production              |
| CI/CD          | GitHub Actions                                    | Lint, typecheck, migrate, deploy   |

### 8.2 Architecture Overview



### 8.3 External Services

| Service                        | Purpose                                        | Phase |
|--------------------------------|------------------------------------------------|-------|
| <b>OpenAI</b>                  | Resume analysis, questions, comment refinement | P0    |
| <b>PostgreSQL (Neon/Azure)</b> | Primary data store                             | P0    |
| <b>S3 / R2 / Azure Blob</b>    | Resume and report storage                      | P1    |
| <b>Resend / SendGrid</b>       | Transactional email                            | P0    |
| <b>Microsoft Entra ID</b>      | Enterprise SSO                                 | P1    |
| <b>Sentry</b>                  | Error monitoring                               | P1    |
| <b>Cloudflare / Vercel</b>     | CDN, TLS, custom domains                       | P1    |

## 8.4 Security Baseline

- Domain-restricted registration with verified email
- Two-factor authentication (TOTP) for users handling candidate PII
- RBAC enforced server-side on every API route
- Resume files: signed URLs, org-scoped paths, encryption at rest
- Audit log for status changes and AI invocations
- AUTH\_SECRET rotation policy
- No PII in AI logs; resume text truncated in prompts (4,000 character cap)

## 9. Phased Delivery Plan

### Phase 0 — Validate Core Loop (4–6 weeks)

**Goal:** TA screens → assigns → interviewer evaluates → recruiter receives PDF → archive

| Deliverable                              | Acceptance Criteria                           |
|------------------------------------------|-----------------------------------------------|
| Wire evaluation comment refinement to UI | Interviewer can refine comments before submit |
| PDF report generation                    | Downloadable PDF from archive                 |
| Candidate entity with basic fields       | Name, email stored with each case             |
| Pipeline in main navigation              | TA sees all candidates by status              |
| Forgot password                          | Email-based reset for web app                 |

### Phase 1 — Recruiter Completeness (6–10 weeks)

| Deliverable                 | Acceptance Criteria                                 |
|-----------------------------|-----------------------------------------------------|
| Mail templates (4 types)    | CRUD, preview, send with variables                  |
| Status-triggered emails     | Correct template fires on reject/shortlist/schedule |
| Cross-project AI suggestion | TA sees alternate project when mismatch detected    |
| Question bank categories    | Project / Role / Experience; reuse without AI       |
| Custom question prompt      | Port from Streamlit legacy                          |
| In-app notifications        | Assignment and submission alerts                    |

### ***Phase 2 — Scheduling & Enterprise (10–16 weeks)***

| Deliverable                       | Acceptance Criteria                                           |
|-----------------------------------|---------------------------------------------------------------|
| Panel availability & slot booking | Interviewer sets windows; TA books slot                       |
| Calendar invites (ICS)            | Candidate receives invite with correct role context           |
| Microsoft Entra SSO               | Production-ready with group → role mapping                    |
| Two-factor authentication         | TOTP enrollment and enforcement                               |
| Admin module v1                   | User invite, role assignment, template management             |
| Reporting dashboard               | Charts: interviews by status, time-to-hire, panel utilization |

### ***Phase 3 — Commercial / Multi-Tenant SaaS (Future)***

- Multi-org per deployment (true tenancy)
- Billing (Stripe)
- Self-service org onboarding
- Custom domains per customer
- API webhooks for ATS integration

## **10. Success Metrics**

| Metric                                       | Baseline            | Target (Post Phase 1) |
|----------------------------------------------|---------------------|-----------------------|
| % interviews resulting in reject after panel | Measure for 30 days | –30%                  |
| Avg. days from resume to decision            | TBD                 | –25%                  |
| Recruiter tool switches per candidate        | Multiple tools      | 1 portal              |

| Metric                      | Baseline    | Target (Post Phase 1)                             |
|-----------------------------|-------------|---------------------------------------------------|
| AI cost per candidate       | TBD         | ≤ 1 analysis + ≤ 1 question gen (with bank reuse) |
| Interviewer prep time       | TBD         | –40% with drafts + question bank                  |
| Template mismatch incidents | Qualitative | Zero reported in pilot                            |

## 11. Risks & Mitigations

| Risk                          | Mitigation                                                 |
|-------------------------------|------------------------------------------------------------|
| AI hallucination in screening | Human-in-the-loop; show resume evidence; never auto-reject |
| Over-reliance on AI questions | Interviewer can edit; bank stores human-approved sets      |
| Email deliverability          | Reputable provider; SPF/DKIM configured                    |
| PII exposure                  | Org-scoped data; encryption; retention policy              |
| Scope creep                   | Strict phase gates; MVP = one complete happy path          |
| Dual codebase confusion       | Deprecate Streamlit after Phase 0 parity                   |

## 12. Team Acceptance Checklist

### ***Problem & Vision***

- Team agrees the problem statement reflects daily pain
- Objectives O1–O6 are accepted as north star
- Non-goals are understood (not building full ATS in MVP)

### ***Scope***

- Phase 0 deliverables are the immediate priority
- Mail templates + PDF are P0 for Phase 0/1
- Panel scheduling deferred to Phase 2 is acceptable
- Admin module deferred to Phase 2 is acceptable

### ***Architecture***

- Next.js + PostgreSQL + OpenAI stack is approved

- Single org per deployment for validation is acceptable
- White-label via env config is sufficient for first external customer

### AI Strategy

- GPT-4o for analysis, GPT-4o-mini for questions approved
- One-time TA analysis reused by interviewers approved
- Question bank reuse to control AI cost approved

### Roles & Users

- TA / Interviewer / Admin RBAC model accepted
- Domain-restricted signup for pilot; SSO in Phase 2

### Open Questions (Resolve in Meeting)

1. **Email provider** — Resend vs. SendGrid vs. corporate SMTP?
2. **Calendar** — Build native scheduling or integrate Outlook first?
3. **2FA method** — TOTP app vs. email OTP?
4. **Pilot size** — How many recruiters and interviewers for Phase 0?
5. **Streamlit sunset** — Target date to stop maintaining legacy app?

### Sign-Off

| Name | Role                       | Date | Approved (Y/N) | Comments |
|------|----------------------------|------|----------------|----------|
|      | TA Lead                    |      |                |          |
|      | Engineering Lead           |      |                |          |
|      | Interviewer Representative |      |                |          |
|      | Product / Management       |      |                |          |

## 13. Appendix — Codebase Mapping

| Feature Area                 | Implementation Location                            |
|------------------------------|----------------------------------------------------|
| Branding / white-label       | web/src/lib/brand.ts, BrandTheme.tsx, .env.example |
| Auth + domain restriction    | web/src/app/login/, web/src/app/register/          |
| Projects / Roles / Questions | web/src/app/setup/, API routes                     |

| Feature Area              | Implementation Location                           |
|---------------------------|---------------------------------------------------|
| AI resume analysis        | web/src/lib/ai/index.ts (analyzeResume, GPT-4o)   |
| Question generation       | web/src/lib/ai/index.ts (GPT-4o-mini)             |
| TA → Interviewer workflow | web/src/lib/db/schema.ts, /assignments, /evaluate |
| Drafts                    | /api/drafts                                       |
| Audit events              | evaluation_events table                           |
| Legacy reference          | Root Streamlit app (pages/, utils/ai_utils.py)    |
| Cloud migration guide     | web/docs/cloud-migration.md                       |

---

\*Document generated for internal team alignment. Version 1.0 — July 2026.\*

**Regenerate PDF:** `python docs/generate-alignment-pdf.py` (requires reportlab)